

Index

Introduction

Nested filters

- Model inference

- Algorithms: Nested particle filter (NPF) and others

Efficient exploration of the parameter space

- Reducing number of particles online

- Numerical results for the Lorenz 63

Conclusions

State-space model

We are interested in systems can be represented by **Markov state-space dynamical models**:

$$\mathbf{x}_t = \mathbf{f}(\mathbf{x}_{t-1}, \boldsymbol{\theta}) + \mathbf{v}_t,$$

$$\mathbf{y}_t = \mathbf{g}(\mathbf{x}_t, \boldsymbol{\theta}) + \mathbf{r}_t,$$

- \mathbf{f} , \mathbf{g} : state transition function and observation function
- \mathbf{v}_t , \mathbf{r}_t : state and observation noises

In terms of a set of **relevant probability density functions (pdfs)**:

- Prior pdfs: $\boldsymbol{\theta} \sim p(\boldsymbol{\theta})$ and $\mathbf{x}_0 \sim p(\mathbf{x}_0)$
- Transition pdf of the state: $\mathbf{x}_t \sim p(\mathbf{x}_t | \mathbf{x}_{t-1}, \boldsymbol{\theta})$
- Conditional pdf of the observation: $\mathbf{y}_t \sim p(\mathbf{y}_t | \mathbf{x}_t, \boldsymbol{\theta})$

State-space model

We are interested in systems can be represented by **Markov state-space dynamical models**:

$$\mathbf{x}_t = \mathbf{f}(\mathbf{x}_{t-1}, \boldsymbol{\theta}) + \mathbf{v}_t,$$

$$\mathbf{y}_t = \mathbf{g}(\mathbf{x}_t, \boldsymbol{\theta}) + \mathbf{r}_t,$$

- \mathbf{f} , \mathbf{g} : state transition function and observation function
- \mathbf{v}_t , \mathbf{r}_t : state and observation noises

In terms of a set of **relevant probability density functions (pdfs)**:

- **Prior pdfs**: $\boldsymbol{\theta} \sim p(\boldsymbol{\theta})$ and $\mathbf{x}_0 \sim p(\mathbf{x}_0)$
- **Transition pdf of the state**: $\mathbf{x}_t \sim p(\mathbf{x}_t | \mathbf{x}_{t-1}, \boldsymbol{\theta})$
- **Conditional pdf of the observation**: $\mathbf{y}_t \sim p(\mathbf{y}_t | \mathbf{x}_t, \boldsymbol{\theta})$

State estimation

Classical filtering methods:

Bayesian estimation of the state variables, $p(\mathbf{x}_t|\mathbf{y}_{1:t}, \boldsymbol{\theta}^*)$, assuming $\boldsymbol{\theta}^*$ is known.

Every time step t :

1. Predictive distribution:

$$p(\mathbf{x}_t|\mathbf{y}_{1:t-1}, \boldsymbol{\theta}^*) = \int p(\mathbf{x}_t|\mathbf{x}_{t-1}, \boldsymbol{\theta}^*)p(\mathbf{x}_{t-1}|\mathbf{y}_{1:t-1}, \boldsymbol{\theta}^*)d\mathbf{x}_{t-1} \quad (1)$$

2. Likelihood: $p(\mathbf{y}_t|\mathbf{x}_t, \boldsymbol{\theta}^*)$

3. Posterior/filtering distribution:

$$p(\mathbf{x}_t|\mathbf{y}_{1:t}, \boldsymbol{\theta}^*) \propto p(\mathbf{y}_t|\mathbf{x}_t, \boldsymbol{\theta}^*)p(\mathbf{x}_t|\mathbf{y}_{1:t-1}, \boldsymbol{\theta}^*) \quad (2)$$

In practice, $\boldsymbol{\theta}^*$ is not known. It is needed to estimate both $\boldsymbol{\theta}$ and \mathbf{x}_t , i.e., $p(\mathbf{x}_t, \boldsymbol{\theta}|\mathbf{y}_{1:t})$.

State estimation

Classical filtering methods:

Bayesian estimation of the state variables, $p(\mathbf{x}_t|\mathbf{y}_{1:t}, \boldsymbol{\theta}^*)$, assuming $\boldsymbol{\theta}^*$ is known.

Every time step t :

1. Predictive distribution:

$$p(\mathbf{x}_t|\mathbf{y}_{1:t-1}, \boldsymbol{\theta}^*) = \int p(\mathbf{x}_t|\mathbf{x}_{t-1}, \boldsymbol{\theta}^*)p(\mathbf{x}_{t-1}|\mathbf{y}_{1:t-1}, \boldsymbol{\theta}^*)d\mathbf{x}_{t-1} \quad (1)$$

2. Likelihood: $p(\mathbf{y}_t|\mathbf{x}_t, \boldsymbol{\theta}^*)$

3. Posterior/filtering distribution:

$$p(\mathbf{x}_t|\mathbf{y}_{1:t}, \boldsymbol{\theta}^*) \propto p(\mathbf{y}_t|\mathbf{x}_t, \boldsymbol{\theta}^*)p(\mathbf{x}_t|\mathbf{y}_{1:t-1}, \boldsymbol{\theta}^*) \quad (2)$$

In practice, $\boldsymbol{\theta}^*$ is not known. It is needed to estimate both $\boldsymbol{\theta}$ and \mathbf{x}_t , i.e., $p(\mathbf{x}_t, \boldsymbol{\theta}|\mathbf{y}_{1:t})$.

State estimation

Classical filtering methods:

Bayesian estimation of the state variables, $p(\mathbf{x}_t|\mathbf{y}_{1:t}, \boldsymbol{\theta}^*)$, assuming $\boldsymbol{\theta}^*$ is known.

Every time step t :

1. Predictive distribution:

$$p(\mathbf{x}_t|\mathbf{y}_{1:t-1}, \boldsymbol{\theta}^*) = \int p(\mathbf{x}_t|\mathbf{x}_{t-1}, \boldsymbol{\theta}^*)p(\mathbf{x}_{t-1}|\mathbf{y}_{1:t-1}, \boldsymbol{\theta}^*)d\mathbf{x}_{t-1} \quad (1)$$

2. Likelihood: $p(\mathbf{y}_t|\mathbf{x}_t, \boldsymbol{\theta}^*)$
3. Posterior/filtering distribution:

$$p(\mathbf{x}_t|\mathbf{y}_{1:t}, \boldsymbol{\theta}^*) \propto p(\mathbf{y}_t|\mathbf{x}_t, \boldsymbol{\theta}^*)p(\mathbf{x}_t|\mathbf{y}_{1:t-1}, \boldsymbol{\theta}^*) \quad (2)$$

In practice, $\boldsymbol{\theta}^*$ is not known. It is needed to estimate both $\boldsymbol{\theta}$ and \mathbf{x}_t , i.e., $p(\mathbf{x}_t, \boldsymbol{\theta}|\mathbf{y}_{1:t})$.

State estimation

Classical filtering methods:

Bayesian estimation of the state variables, $p(\mathbf{x}_t|\mathbf{y}_{1:t}, \boldsymbol{\theta}^*)$, assuming $\boldsymbol{\theta}^*$ is known.

Every time step t :

1. Predictive distribution:

$$p(\mathbf{x}_t|\mathbf{y}_{1:t-1}, \boldsymbol{\theta}^*) = \int p(\mathbf{x}_t|\mathbf{x}_{t-1}, \boldsymbol{\theta}^*)p(\mathbf{x}_{t-1}|\mathbf{y}_{1:t-1}, \boldsymbol{\theta}^*)d\mathbf{x}_{t-1} \quad (1)$$

2. Likelihood: $p(\mathbf{y}_t|\mathbf{x}_t, \boldsymbol{\theta}^*)$
3. Posterior/filtering distribution:

$$p(\mathbf{x}_t|\mathbf{y}_{1:t}, \boldsymbol{\theta}^*) \propto p(\mathbf{y}_t|\mathbf{x}_t, \boldsymbol{\theta}^*)p(\mathbf{x}_t|\mathbf{y}_{1:t-1}, \boldsymbol{\theta}^*) \quad (2)$$

In practice, $\boldsymbol{\theta}^*$ is not known. It is needed to estimate both $\boldsymbol{\theta}$ and \mathbf{x}_t , i.e., $p(\mathbf{x}_t, \boldsymbol{\theta}|\mathbf{y}_{1:t})$.

State-of-the-art methods

Methods for Bayesian inference of both θ and \mathbf{x}_t :

- **particle Markov chain Monte Carlo (PMCMC)**¹
- **sequential Monte Carlo square (SMC²)**²
- **nested particle filters (NPFs)**³

- They can **quantify the uncertainty** or estimation error.
- They can be applied to a **broad class of models**.
- They provide **theoretical guarantees**.
- Both PMCMC and SMC² are **batch techniques**, while the NPF is a **recursive method**.

¹Andrieu, Doucet, and Holenstein 2010.

²Chopin, Jacob, and Papaspiliopoulos 2013.

³Crisan and Míguez 2018.

State-of-the-art methods

Methods for Bayesian inference of both θ and \mathbf{x}_t :

- **particle Markov chain Monte Carlo (PMCMC)**¹
- **sequential Monte Carlo square (SMC²)**²
- **nested particle filters (NPFs)**³
 - They can **quantify the uncertainty** or estimation error.
 - They can be applied to a **broad class of models**.
 - They provide **theoretical guarantees**.
 - Both PMCMC and SMC² are **batch techniques**, while the NPF is a **recursive method**.

¹Andrieu, Doucet, and Holenstein 2010.

²Chopin, Jacob, and Papaspiliopoulos 2013.

³Crisan and Míguez 2018.

State-of-the-art methods

Methods for Bayesian inference of both θ and \mathbf{x}_t :

- **particle Markov chain Monte Carlo (PMCMC)**¹
- **sequential Monte Carlo square (SMC²)**²
- **nested particle filters (NPFs)**³
 - They can **quantify the uncertainty** or estimation error.
 - They can be applied to a **broad class of models**.
 - They provide **theoretical guarantees**.
 - Both PMCMC and SMC² are **batch techniques**, while the NPF is a **recursive method**.

¹Andrieu, Doucet, and Holenstein 2010.

²Chopin, Jacob, and Papaspiliopoulos 2013.

³Crisan and Míguez 2018.

Index

Introduction

Nested filters

Model inference

Algorithms: Nested particle filter (NPF) and others

Efficient exploration of the parameter space

Reducing number of particles online

Numerical results for the Lorenz 63

Conclusions

Model inference

We aim at computing the **joint posterior pdf** $p(\boldsymbol{\theta}, \mathbf{x}_t | \mathbf{y}_{1:t})$, that can be written as

$$p(\mathbf{x}_t, \boldsymbol{\theta} | \mathbf{y}_{1:t}) = \underbrace{p(\mathbf{x}_t | \boldsymbol{\theta}, \mathbf{y}_{1:t})}_{2^{nd} \text{ layer}} \underbrace{p(\boldsymbol{\theta} | \mathbf{y}_{1:t})}_{1^{st} \text{ layer}}$$

→ The **key difficulty** in this class of models is **the Bayesian estimation of the parameter vector $\boldsymbol{\theta}$** .

Model inference

At every time step t :

$$\underbrace{p(\theta|y_{1:t-1})}_{\text{Pred. pdf of } \theta}$$

1st layer

Filtering (given θ)

$$\underbrace{p(x_t|\theta, y_{1:t-1})}_{\text{Pred. pdf of } x}$$

$$p(y_t|x_t, \theta)$$

$$\underbrace{p(x_t|\theta, y_{1:t})}_{\text{Post. pdf of } x}$$

2nd layer

$$p(y_t|\theta, y_{1:t-1})$$

$$\underbrace{p(\theta|y_{1:t})}_{\text{Post. pdf of } \theta} \propto p(y_t|\theta, y_{1:t-1})p(\theta|y_{1:t-1})$$

Model inference

At every time step t :

$$\underbrace{p(\theta|y_{1:t-1})}_{\text{Pred. pdf of } \theta}$$

1st layer

Filtering (given θ)

$$\underbrace{p(x_t|\theta, y_{1:t-1})}_{\text{Pred. pdf of } x}$$

$$p(y_t|x_t, \theta)$$

$$\underbrace{p(x_t|\theta, y_{1:t})}_{\text{Post. pdf of } x}$$

2nd layer

$$p(y_t|\theta, y_{1:t-1}) = \int p(y_t|x_t, \theta) p(x_t|\theta, y_{1:t-1}) dx_t$$

$$\underbrace{p(\theta|y_{1:t})}_{\text{Post. pdf of } \theta} \propto p(y_t|\theta, y_{1:t-1}) p(\theta|y_{1:t-1})$$

Model inference

At every time step t :

$$\underbrace{p(\theta|y_{1:t-1})}_{\text{Pred. pdf of } \theta}$$

1st layer

Filtering (given θ)

$$\underbrace{p(x_t|\theta, y_{1:t-1})}_{\text{Pred. pdf of } x}$$

$$p(y_t|x_t, \theta)$$

$$\underbrace{p(x_t|\theta, y_{1:t})}_{\text{Post. pdf of } x}$$

2nd layer

$$p(y_t|\theta, y_{1:t-1}) = \int p(y_t|x_t, \theta) p(x_t|\theta, y_{1:t-1}) dx_t$$

$$\underbrace{p(\theta|y_{1:t})}_{\text{Post. pdf of } \theta} \propto p(y_t|\theta, y_{1:t-1}) p(\theta|y_{1:t-1})$$

Model inference

At every time step t :

$$\underbrace{p(\theta|\mathbf{y}_{1:t-1})}_{\text{Pred. pdf of } \theta}$$

1st layer

Filtering (given θ)

$$\underbrace{p(\mathbf{x}_t|\theta, \mathbf{y}_{1:t-1})}_{\text{Pred. pdf of } \mathbf{x}} = \int p(\mathbf{x}_t|\mathbf{x}_{t-1}, \theta) p(\mathbf{x}_{t-1}|\theta, \mathbf{y}_{1:t-1}) d\mathbf{x}_{t-1}$$

$$p(\mathbf{y}_t|\mathbf{x}_t, \theta)$$

$$\underbrace{p(\mathbf{x}_t|\theta, \mathbf{y}_{1:t})}_{\text{Post. pdf of } \mathbf{x}} \propto p(\mathbf{y}_t|\mathbf{x}_t, \theta) p(\mathbf{x}_t|\theta, \mathbf{y}_{1:t-1})$$

2nd layer

$$p(\mathbf{y}_t|\theta, \mathbf{y}_{1:t-1}) = \int p(\mathbf{y}_t|\mathbf{x}_t, \theta) p(\mathbf{x}_t|\theta, \mathbf{y}_{1:t-1}) d\mathbf{x}_t$$

$$\underbrace{p(\theta|\mathbf{y}_{1:t})}_{\text{Post. pdf of } \theta} \propto p(\mathbf{y}_t|\theta, \mathbf{y}_{1:t-1}) p(\theta|\mathbf{y}_{1:t-1})$$

Model inference

At every time step t :

$$\underbrace{p(\theta|\mathbf{y}_{1:t-1})}_{\text{Pred. pdf of } \theta}$$

1st layer

Filtering (given θ)

$$\underbrace{p(\mathbf{x}_t|\theta, \mathbf{y}_{1:t-1})}_{\text{Pred. pdf of } \mathbf{x}} = \int p(\mathbf{x}_t|\mathbf{x}_{t-1}, \theta) p(\mathbf{x}_{t-1}|\theta, \mathbf{y}_{1:t-1}) d\mathbf{x}_{t-1}$$

$$p(\mathbf{y}_t|\mathbf{x}_t, \theta)$$

$$\underbrace{p(\mathbf{x}_t|\theta, \mathbf{y}_{1:t})}_{\text{Post. pdf of } \mathbf{x}} \propto p(\mathbf{y}_t|\mathbf{x}_t, \theta) p(\mathbf{x}_t|\theta, \mathbf{y}_{1:t-1})$$

2nd layer

$$p(\mathbf{y}_t|\theta, \mathbf{y}_{1:t-1}) = \int p(\mathbf{y}_t|\mathbf{x}_t, \theta) p(\mathbf{x}_t|\theta, \mathbf{y}_{1:t-1}) d\mathbf{x}_t$$

$$\underbrace{p(\theta|\mathbf{y}_{1:t})}_{\text{Post. pdf of } \theta} \propto p(\mathbf{y}_t|\theta, \mathbf{y}_{1:t-1}) p(\theta|\mathbf{y}_{1:t-1})$$

Naive importance sampling approximation

Initialisation: Draw $\{\theta^i\}_{i=1}^{N_\theta}$ from $p(\theta)$

At $t \geq 1$ and for every θ^i , $i = 1, \dots, N_\theta$:

SMC (N_θ samples)
to approximate $p(\theta|\mathbf{y}_{1:t})$

For $j = 1, \dots, N_x$:

SMC (N_x samples)
to approximate $p(\mathbf{y}_t | \mathbf{y}_{1:t-1}, \theta^i)$

Naive importance sampling approximation

Initialisation: Draw $\{\theta^i\}_{i=1}^{N_\theta}$ from $p(\theta)$

At $t \geq 1$ and for every θ^i , $i = 1, \dots, N_\theta$:

SMC (N_θ samples)
to approximate $p(\theta|\mathbf{y}_{1:t})$

For $j = 1, \dots, N_x$:

SMC (N_x samples)
to approximate $p(\mathbf{y}_t | \mathbf{y}_{1:t-1}, \theta^j)$

- Draw $\tilde{\mathbf{x}}_t^{i,j} \sim p(\mathbf{x}_t | \theta^i, \mathbf{y}_{1:t-1})$
- Weights: $\tilde{u}_t^{i,j} \propto p(\mathbf{y}_t | \tilde{\mathbf{x}}_t^{i,j}, \theta^i)$
- Resampling: for $m = 1, \dots, N_x$, $\tilde{\mathbf{x}}_t^{i,j} = \tilde{\mathbf{x}}_t^{i,m}$
with prob. $u_t^{i,m} = \frac{\tilde{u}_t^{i,m}}{\sum_{j=1}^{N_x} \tilde{u}_t^{i,j}}$

Naive importance sampling approximation

Initialisation: Draw $\{\theta^i\}_{i=1}^{N_\theta}$ from $p(\theta)$

At $t \geq 1$ and for every θ^i , $i = 1, \dots, N_\theta$:

SMC (N_θ samples)
to approximate $p(\theta|\mathbf{y}_{1:t})$

For $j = 1, \dots, N_x$:

SMC (N_x samples)
to approximate $p(\mathbf{y}_t | \mathbf{y}_{1:t-1}, \theta^i)$

- Draw $\bar{\mathbf{x}}_t^{i,j} \sim p(\mathbf{x}_t | \boldsymbol{\theta}^i, \mathbf{y}_{1:t-1})$
- Weights: $\tilde{u}_t^{i,j} \propto p(\mathbf{y}_t | \bar{\mathbf{x}}_t^{i,j}, \boldsymbol{\theta}^i)$
- Resampling: for $m = 1, \dots, N_x$, $\tilde{\mathbf{x}}_t^{i,j} = \bar{\mathbf{x}}_t^{i,m}$
with prob. $u_t^{i,m} = \frac{\tilde{u}_t^{i,m}}{\sum_{j=1}^{N_x} \tilde{u}_t^{i,j}}$

Naive importance sampling approximation

Initialisation: Draw $\{\theta^i\}_{i=1}^{N_\theta}$ from $p(\theta)$

At $t \geq 1$ and for every θ^i , $i = 1, \dots, N_\theta$:

SMC (N_θ samples)
to approximate $p(\theta|\mathbf{y}_{1:t})$

For $j = 1, \dots, N_x$:

SMC (N_x samples)
to approximate $p(\mathbf{y}_t | \mathbf{y}_{1:t-1}, \theta^i)$

- Draw $\bar{\mathbf{x}}_t^{i,j} \sim p(\mathbf{x}_t | \boldsymbol{\theta}^i, \mathbf{y}_{1:t-1})$
- Weights: $\tilde{u}_t^{i,j} \propto p(\mathbf{y}_t | \bar{\mathbf{x}}_t^{i,j}, \boldsymbol{\theta}^i)$
- Resampling: for $m = 1, \dots, N_x$, $\tilde{\mathbf{x}}_t^{i,j} = \bar{\mathbf{x}}_t^{i,m}$
with prob. $u_t^{i,m} = \frac{\tilde{u}_t^{i,m}}{\sum_{j=1}^{N_x} \tilde{u}_t^{i,j}}$

Naive importance sampling approximation

- Careful with $p(\theta)$: after several time steps the filter **degenerates**
- Possible solution: drawing $\{\theta_t^i\} \sim p(\theta | \mathbf{y}_{1:t-1})$ at each time step \rightarrow **re-running from scratch the filter for \mathbf{x}** (i.e., not recursive anymore)

- NPF \rightarrow **jittering**: $\bar{\theta}_t^i \sim \kappa_{N_\theta}(d\theta|\theta')$, where

$$\kappa_{N_\theta}(d\theta|\theta') = (1 - \epsilon_{N_\theta})\delta_{\theta'}(\theta) + \epsilon_{N_\theta}\kappa(d\theta|\theta')$$

- $0 < \epsilon_{N_\theta} \leq \frac{1}{\sqrt{N_\theta}}$
- $\kappa(d\theta|\theta')$ is an arbitrary Markov kernel with mean θ' and finite variance, e.g., $\kappa(d\theta|\theta') = \mathcal{N}(\theta|\theta', \tilde{\sigma}^2 I)$, with $\tilde{\sigma}^2 < \infty$.
- Guarantees convergence to the true posterior when $N \rightarrow \infty$

Naive importance sampling approximation

- Careful with $p(\theta)$: after several time steps the filter **degenerates**
- Possible solution: drawing $\{\theta_t^i\} \sim p(\theta | \mathbf{y}_{1:t-1})$ at each time step \rightarrow **re-running from scratch the filter for \mathbf{x}** (i.e., not recursive anymore)

- NPF \rightarrow **jittering**: $\bar{\theta}_t^i \sim \kappa_{N_\theta}(d\theta | \theta')$, where

$$\kappa_{N_\theta}(d\theta | \theta') = (1 - \epsilon_{N_\theta})\delta_{\theta'}(\theta) + \epsilon_{N_\theta}\kappa(d\theta | \theta')$$

- $0 < \epsilon_{N_\theta} \leq \frac{1}{\sqrt{N_\theta}}$
- $\kappa(d\theta | \theta')$ is an arbitrary Markov kernel with mean θ' and finite variance, e.g., $\kappa(d\theta | \theta') = \mathcal{N}(\theta | \theta', \tilde{\sigma}^2 \mathbf{I})$, with $\tilde{\sigma}^2 < \infty$.
- Guarantees convergence to the true posterior when $N \rightarrow \infty$

Nested particle filter (NPF)⁴For $i = 1, \dots, N_\theta$:

- **Jittering**: Draw $\bar{\theta}_t^i \sim \kappa_{N_\theta}(d\theta | \theta_{t-1}^i)$

SMC (N_θ samples)
to approximate $p(\theta | \mathbf{y}_{1:t})$ Given $\bar{\theta}_t^i$, for $j = 1, \dots, N_x$:

- Draw $\bar{\mathbf{x}}_t^{i,j} \sim p(\mathbf{x}_t | \bar{\theta}_t^i, \mathbf{y}_{1:t-1})$

SMC (N_x samples)
to approximate $p(\mathbf{y}_t | \mathbf{y}_{1:t-1}, \bar{\theta}_t^i)$

- Weights: $\tilde{u}_t^{i,j} \propto p(\mathbf{y}_t | \bar{\mathbf{x}}_t^{i,j}, \bar{\theta}_t^i)$

- Resampling: for $m = 1, \dots, N_x$, $\tilde{\mathbf{x}}_t^{i,j} = \bar{\mathbf{x}}_t^{i,m}$
with prob. $u_t^{i,m} = \frac{\tilde{u}_t^{i,m}}{\sum_{j=1}^{N_x} \tilde{u}_t^{i,j}}$

- **Likelihood of $\bar{\theta}_t^i$** : $\tilde{w}_t^i = \frac{1}{N_x} \sum_{j=1}^{N_x} \tilde{u}_t^{i,j}$
- **Resampling**: for $l = 1, \dots, N_\theta$, $\{\theta_t^l, \{\mathbf{x}_t^{l,j}\}_{1 \leq j \leq N_x}\} = \{\bar{\theta}_t^l, \{\tilde{\mathbf{x}}_t^{l,j}\}_{1 \leq j \leq N_x}\}$
with prob. w_t^l , so that $p(\theta | \mathbf{y}_{1:t}) = \frac{1}{N_\theta} \sum_{i=1}^{N_\theta} \delta_{\theta_t^i}(d\theta)$

⁴Crisan and Miguez 2017.

Family of nested filters

1. Nested particle filters (NPFs)⁵.

- Both layers → Sequential Monte Carlo (SMC) methods
High computational complexity: $N_\theta \times N_x$

2. Nested hybrid filters (NHF)⁶.

- θ -layer → Monte Carlo-based methods (e.g., SMC or SQMC)
- x -layer → Gaussian techniques (e.g., EKF or EnKF)

3. Nested Gaussian filters (NGF)⁷.

- θ -layer → Deterministic sampling methods (e.g., UKF).
- x -layer → Gaussian techniques (e.g., EKF or EnKF).

⁵Crisan and Míguez 2018.

⁶Pérez-Vieites, Mariño, and Míguez 2017.

⁷Pérez-Vieites and Míguez 2021.

Family of nested filters

1. Nested particle filters (NPFs)⁵.

- Both layers → Sequential Monte Carlo (SMC) methods
High computational complexity: $N_\theta \times N_x$

2. Nested hybrid filters (NHF)⁶.

- θ -layer → Monte Carlo-based methods (e.g., SMC or SQMC)
- x -layer → Gaussian techniques (e.g., EKF or EnKF)

3. Nested Gaussian filters (NGF)⁷.

- θ -layer → Deterministic sampling methods (e.g., UKF).
- x -layer → Gaussian techniques (e.g., EKF or EnKF).

⁵Crisan and Míguez 2018.

⁶Pérez-Vieites, Mariño, and Míguez 2017.

⁷Pérez-Vieites and Míguez 2021.

Family of nested filters

1. Nested particle filters (NPFs)⁵.

- Both layers → Sequential Monte Carlo (SMC) methods
High computational complexity: $N_\theta \times N_x$

2. Nested hybrid filters (NHF)⁶.

- θ -layer → Monte Carlo-based methods (e.g., SMC or SQMC)
- x -layer → Gaussian techniques (e.g., EKF or EnKF)

3. Nested Gaussian filters (NGF)⁷.

- θ -layer → Deterministic sampling methods (e.g., UKF).
- x -layer → Gaussian techniques (e.g., EKF or EnKF).

⁵Crisan and Míguez 2018.

⁶Pérez-Vieites, Mariño, and Míguez 2017.

⁷Pérez-Vieites and Míguez 2021.

Reducing number of particles online

Problem: Great amount of samples ($N_\theta \times N_x$) and **waste of computational effort** when they are not well chosen.

Possible approach: reducing automatically the number of samples, N_θ , when the performance is no longer compromised.

We studied the case for a nested Gaussian filter that implements:

- **Quadrature Kalman filter (QKF)** in the θ -layer, with $N_\theta = \alpha^{d_\theta}, \alpha > 1$.

The hyperparameter α will depend on t , so the number of samples is now defined as $N_{\theta,t} = \alpha_t^{d_\theta}$.

Reducing number of particles online

Problem: Great amount of samples ($N_\theta \times N_x$) and **waste of computational effort** when they are not well chosen.

Possible approach: reducing automatically the number of samples, N_θ , when the performance is no longer compromised.

We studied the case for a nested Gaussian filter that implements:

- **Quadrature Kalman filter (QKF) in the θ -layer**, with $N_\theta = \alpha^{d_\theta}, \alpha > 1$.

The hyperparameter α will depend on t , so the number of samples is now defined as $N_{\theta,t} = \alpha_t^{d_\theta}$.

Reducing number of particles online

Problem: Great amount of samples ($N_\theta \times N_x$) and **waste of computational effort** when they are not well chosen.

Possible approach: reducing automatically the number of samples, N_θ , when the performance is no longer compromised.

We studied the case for a nested Gaussian filter that implements:

- **Quadrature Kalman filter (QKF) in the θ -layer**, with $N_\theta = \alpha^{d_\theta}, \alpha > 1$.

The hyperparameter α will depend on t , so the number of samples is now defined as $N_{\theta,t} = \alpha_t^{d_\theta}$.

Reducing number of particles online

Problem: Great amount of samples ($N_\theta \times N_x$) and **waste of computational effort** when they are not well chosen.

Possible approach: reducing automatically the number of samples, N_θ , when the performance is no longer compromised.

We studied the case for a nested Gaussian filter that implements:

- **Quadrature Kalman filter (QKF) in the θ -layer**, with $N_\theta = \alpha^{d_\theta}, \alpha > 1$.

The hyperparameter α will depend on t , so the number of samples is now defined as $N_{\theta,t} = \alpha_t^{d_\theta}$.

Adaptive reduction rule

New statistic to decide when to reduce $N_{\theta,t}$:

$$\rho_t = \frac{1}{\sum_{n=1}^{N_{\theta,t}} (\bar{s}_t^n)^2} \quad \text{with} \quad \bar{s}_t^n = \frac{p(\mathbf{y}_t | \mathbf{y}_{1:t-1}, \boldsymbol{\theta}_t^n)}{\sum_{n=1}^{N_{\theta,t}} p(\mathbf{y}_t | \mathbf{y}_{1:t-1}, \boldsymbol{\theta}_t^n)}$$

The statistic takes

- its **minimum value** in $\rho_t = 1$, which occurs when **only one** $p(\mathbf{y}_t | \mathbf{y}_{1:t-1}, \boldsymbol{\theta}_t^n)$, for $n = 1, \dots, N_{\theta,t}$, is **different from zero**; and
- its **maximum value** in $\rho_t = N_{\theta,t}$, when for **all** $n = 1, \dots, N_{\theta,t}$, the evaluations $p(\mathbf{y}_t | \mathbf{y}_{1:t-1}, \boldsymbol{\theta}_t^n)$ **are equal**.

The adaptive reduction rule:

- If $\frac{\rho_t}{N_{\theta,t}} > 1 - \epsilon$ (ρ_t is close to its maximum value),
$$N_{\theta,t+1} = (\alpha_t - 1)^{d_\theta} < N_{\theta,t}, \text{ with } N_{\theta,t+1} > N_{\min}.$$
- Otherwise, $N_{\theta,t+1} = N_{\theta,t}$.

Adaptive reduction rule

New statistic to decide when to reduce $N_{\theta,t}$:

$$\rho_t = \frac{1}{\sum_{n=1}^{N_{\theta,t}} (\bar{s}_t^n)^2} \quad \text{with} \quad \bar{s}_t^n = \frac{p(\mathbf{y}_t | \mathbf{y}_{1:t-1}, \boldsymbol{\theta}_t^n)}{\sum_{n=1}^{N_{\theta,t}} p(\mathbf{y}_t | \mathbf{y}_{1:t-1}, \boldsymbol{\theta}_t^n)}$$

The statistic takes

- its **minimum value** in $\rho_t = 1$, which occurs when **only one** $p(\mathbf{y}_t | \mathbf{y}_{1:t-1}, \boldsymbol{\theta}_t^n)$, for $n = 1, \dots, N_{\theta,t}$, is **different from zero**; and
- its **maximum value** in $\rho_t = N_{\theta,t}$, when for **all** $n = 1, \dots, N_{\theta,t}$, the evaluations $p(\mathbf{y}_t | \mathbf{y}_{1:t-1}, \boldsymbol{\theta}_t^n)$ **are equal**.

The **adaptive reduction rule**:

- If $\frac{\rho_t}{N_{\theta,t}} > 1 - \epsilon$ (ρ_t is close to its maximum value),
$$N_{\theta,t+1} = (\alpha_t - 1)^{d_\theta} < N_{\theta,t}, \text{ with } N_{\theta,t+1} > N_{\min}.$$
- Otherwise, $N_{\theta,t+1} = N_{\theta,t}$.

Index

Introduction

Nested filters

Model inference

Algorithms: Nested particle filter (NPF) and others

Efficient exploration of the parameter space

Reducing number of particles online

Numerical results for the Lorenz 63

Conclusions

Numerical results

- Applying a discretization method with step Δ , we obtain

$$x_{1,t+1} = x_{1,t} - \Delta S(x_{1,t} - x_{2,t}) + \sqrt{\Delta} \sigma v_{1,t},$$

$$x_{2,t+1} = x_{2,t} + \Delta[(R - x_{3,t})x_{1,t} - x_{2,t}] + \sqrt{\Delta} \sigma v_{2,t},$$

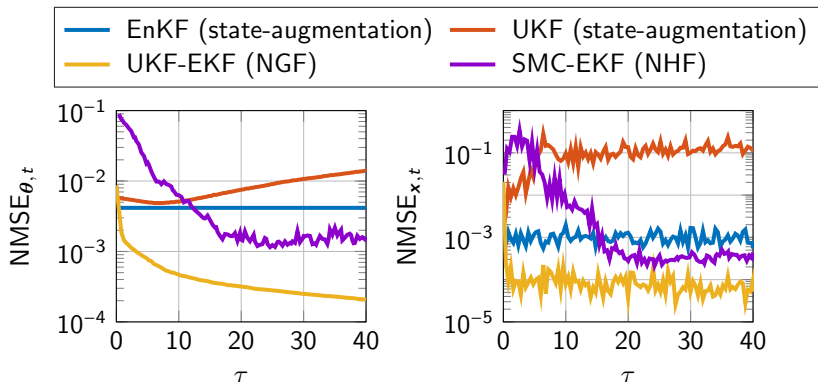
$$x_{3,t+1} = x_{3,t} + \Delta(x_{1,t}x_{2,t} - Bx_{3,t}) + \sqrt{\Delta} \sigma v_{3,t},$$

- We assume linear observations of the form

$$\mathbf{y}_t = k_o \begin{bmatrix} x_{1,t} \\ x_{3,t} \end{bmatrix} + \mathbf{r}_t,$$

where k_o is a fixed known parameter and $\mathbf{r}_t \sim \mathcal{N}(\mathbf{r}_t | \mathbf{0}, \sigma_y^2 \mathbf{I}_2)$.

Numerical results⁸

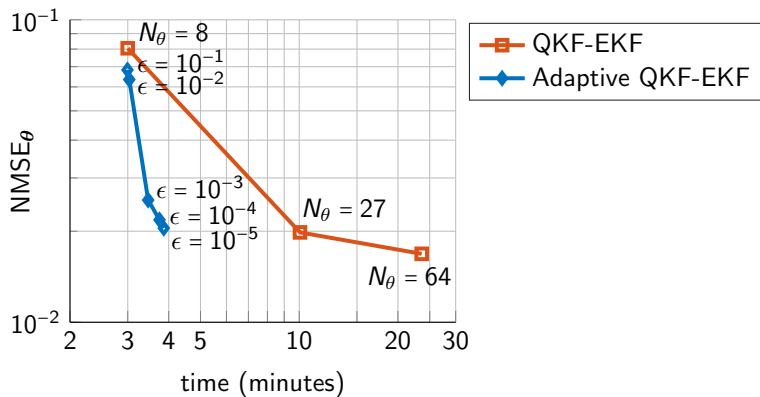


→ The nested schemes outperform the augmented-state methods.

→ The UKF-EKF is three times faster than SMC-EKF.

⁸Pérez-Vieites and Míguez 2021.

Numerical results⁹



1. **QKF-EKF** for different fixed $N_{\theta} = \{8, 27, 64\}$.
2. **Adaptive QKF-EKF** with $N_{\theta,1} = 64$.

Conclusions

1. The nested methodology is **online and flexible**. It admits different types of filtering techniques in each layer, leading to a **set of algorithms**.
2. For a **further reduction of the computational complexity**. Automatic reduction of N_θ when points become less informative → reduction of cost for a given performance.

Conclusions

1. The nested methodology is **online and flexible**. It admits different types of filtering techniques in each layer, leading to a **set of algorithms**.
2. For a **further reduction of the computational complexity**. Automatic reduction of N_θ when points become less informative → reduction of cost for a given performance.

